

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import os

BASE_dir = os.getcwd()
DATA_PATH = os.path.abspath(BASE_dir + '..\data\smart_home_energy.csv')
FIG_PATH = os.path.abspath(BASE_dir + '../reports/Milestone1/figures')

df_initial = pd.read_csv(DATA_PATH)
df_initial
```

Out [1]:

	home_id	timestamp	device_id	device_type	room	status	power_watt	user_present	activity	indoor_temp	outdoor_temp
0	1	2022-01-01 00:00:00	air_conditioner1	air_conditioner	bedroom	off	0.000000	1	sleeping	11.4	11.9
1	1	2022-01-01 00:00:00	light1	light	living_room	on	105.880000	1	sleeping	11.4	11.9
2	1	2022-01-01 00:00:00	tv1	tv	living_room	off	0.000000	1	sleeping	11.4	11.9
3	1	2022-01-01 00:00:00	fridge1	fridge	kitchen	on	223.460000	1	sleeping	11.4	11.9
4	1	2022-01-01 00:00:00	washer1	washer	laundry_room	off	0.000000	1	sleeping	11.4	11.9
...
1751995	10	2022-12-31 23:45:00	air_conditioner10	air_conditioner	bedroom	off	0.000000	1	sleeping	10.8	11.1
1751996	10	2022-12-31 23:45:00	light10	light	living_room	off	0.000000	1	sleeping	10.8	11.1
1751997	10	2022-12-31 23:45:00	tv10	tv	living_room	off	0.000000	1	sleeping	10.8	11.1
1751998	10	2022-12-31 23:45:00	fridge10	fridge	kitchen	on	261.350000	1	sleeping	10.8	11.1
1751999	10	2022-12-31 23:45:00	washer10	washer	laundry_room	on	1884.819597	1	sleeping	10.8	11.1

1752000 rows × 16 columns

In [2]: `print(df_initial.columns)`

```
Index(['home_id', 'timestamp', 'device_id', 'device_type', 'room', 'status',
       'power_watt', 'user_present', 'activity', 'indoor_temp', 'outdoor_temp',
       'humidity', 'light_level', 'day_of_week', 'hour_of_day', 'price_kWh'],
      dtype='object')
```

Checking the data types for all columns

In [3]: `df_initial.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1752000 entries, 0 to 1751999
Data columns (total 16 columns):
 #   Column        Dtype  
 --- 
 0   home_id       int64  
 1   timestamp     object  
 2   device_id     object  
 3   device_type   object  
 4   room          object  
 5   status         object  
 6   power_watt    float64 
 7   user_present  int64  
 8   activity       object  
 9   indoor_temp   float64 
 10  outdoor_temp  float64 
 11  humidity      float64 
 12  light_level   float64 
 13  day_of_week   int64  
 14  hour_of_day   int64  
 15  price_kWh    int64  
dtypes: float64(5), int64(5), object(6)
memory usage: 213.9+ MB
```

Converting the timestamp column to appropriate type for further analysis

```
In [4]: df_initial['timestamp'] = pd.to_datetime(df_initial['timestamp'])
df_initial.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1752000 entries, 0 to 1751999
Data columns (total 16 columns):
 #   Column        Dtype  
--- 
 0   home_id       int64  
 1   timestamp     datetime64[ns]
 2   device_id     object  
 3   device_type   object  
 4   room          object  
 5   status         object  
 6   power_watt    float64 
 7   user_present  int64  
 8   activity       object  
 9   indoor_temp   float64 
10  outdoor_temp  float64 
11  humidity       float64 
12  light_level   float64 
13  day_of_week   int64  
14  hour_of_day   int64  
15  price_kWh    int64  
dtypes: datetime64[ns](1), float64(5), int64(5), object(5)
memory usage: 213.9+ MB
```

Checking for null values column-wise

```
In [5]: for attr in df_initial.columns:
    print(df_initial[attr].isnull().value_counts())
    print()
```

```
home_id
False    1752000
Name: count, dtype: int64

timestamp
False    1752000
Name: count, dtype: int64

device_id
False    1752000
Name: count, dtype: int64

device_type
False    1752000
Name: count, dtype: int64

room
False    1752000
Name: count, dtype: int64

status
False    1752000
Name: count, dtype: int64

power_watt
False    1752000
Name: count, dtype: int64

user_present
False    1752000
Name: count, dtype: int64

activity
False    1752000
Name: count, dtype: int64

indoor_temp
False    1752000
Name: count, dtype: int64

outdoor_temp
False    1752000
Name: count, dtype: int64

humidity
False    1752000
Name: count, dtype: int64

light_level
False    1752000
Name: count, dtype: int64

day_of_week
False    1752000
Name: count, dtype: int64

hour_of_day
False    1752000
Name: count, dtype: int64

price_kWh
False    1752000
Name: count, dtype: int64
```

No null values found

Since all columns have repeated valid data, checking for duplicates isn't necessary

Type of appliances considered in datasets

```
In [6]: df_initial['device_type'].value_counts()
```

```
Out [6]: device_type
air_conditioner    350400
light             350400
tv                350400
fridge            350400
washer            350400
Name: count, dtype: int64
```

Finding the difference in intervals between timestamps

```
In [7]: df_unique_ts = df_initial.drop_duplicates(subset='timestamp')['timestamp']
df_unique_ts.diff().value_counts().head()
```

```
Out [7]: timestamp
0 days 00:15:00    35039
Name: count, dtype: int64
```

Since there is 15 min (15min) difference between each timestamp, checking for any missing interval

```
In [8]: full_range = pd.date_range(start=df_unique_ts.min(), end=df_unique_ts.max(), freq='15min')
missing = full_range.difference(df_unique_ts)
print(f'missing timestamps: {missing}')
```



```
missing timestamps: DatetimeIndex([], dtype='datetime64[ns]', freq='15min')
```

no missing timestamp

Computing Energy in kWh from power logs for each interval

```
In [9]: df_initial['energy_kWh'] = (df_initial['power_watt'] * 0.25) / 1000 # since power is energy consumed in unit
df_initial['energy_kWh']
```

```
Out [9]: 0      0.000000
1      0.026470
2      0.000000
3      0.055865
4      0.000000
...
1751995 0.000000
1751996 0.000000
1751997 0.000000
1751998 0.065338
1751999 0.471205
Name: energy_kWh, Length: 1752000, dtype: float64
```

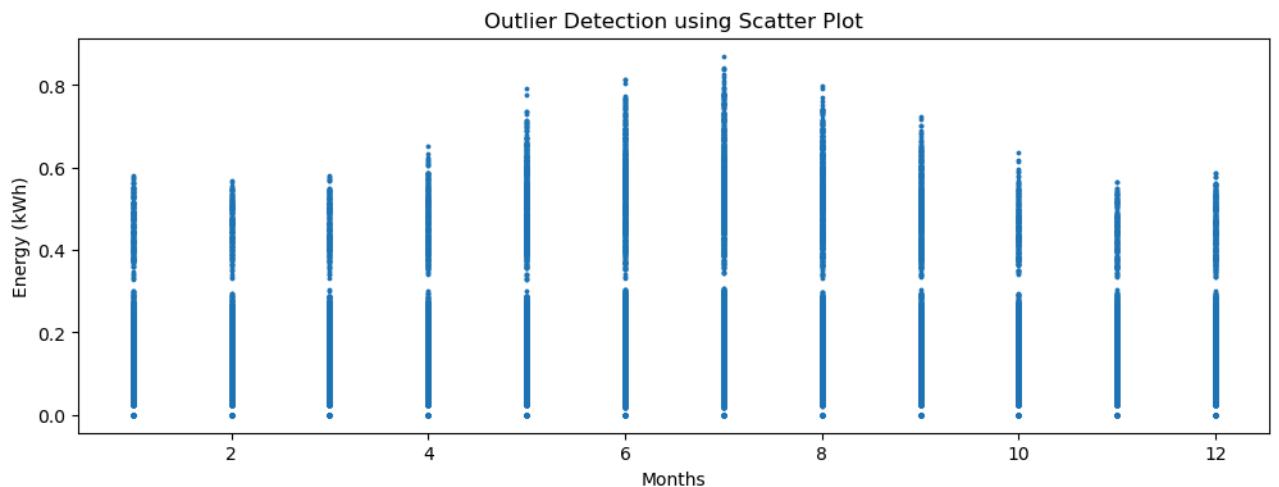
Checking for any invalid values of energy consumption (negative)

```
In [10]: (df_initial['energy_kWh'] < 0).value_counts() # as no device can consume negative energy
```

```
Out [10]: energy_kWh
False    1752000
Name: count, dtype: int64
```

Outlier Detection using scatter plot in month-wise distribution

```
In [11]: plt.figure(figsize=(12,4))
plt.scatter(df_initial['timestamp'].dt.month, df_initial['energy_kWh'], s = 3)
plt.title("Outlier Detection using Scatter Plot")
plt.xlabel('Months')
plt.ylabel('Energy (kWh)')
plt.savefig(FIG_PATH+'\Energy_Month_Wise.png')
plt.show()
```



no outlier observed


Application-wise consumption distribution wrt different features

```
In [12]: df_subset1 = df_initial.pivot_table(index='device_type', values='energy_kWh', aggfunc='sum')
df_subset1.round(2)
```

```
Out [12]: energy_kWh
```

device_type	energy_kWh
air_conditioner	13001.83
fridge	23248.45
light	3539.48
tv	7448.75
washer	5315.65

```
In [13]: df_subset1.plot(kind='bar')
plt.title('Energy Consumption per appliance in a year')
plt.xlabel('Device Type')
plt.ylabel('Energy in kWh')
plt.savefig(FIG_PATH+'\Energy_Apppliance_Wise.png')
plt.show()
```




```
In [14]: df_subset2 = df_initial.pivot_table(index='activity', columns='device_type', values='energy_kWh', aggfunc='sum')
df_subset2.round(2)
```

```
Out [14]: device_type  air_conditioner    fridge     light       tv      washer
activity
away        803.07    4821.68   153.43   350.79   700.70
```

device_type	air_conditioner	fridge	light	tv	washer
activity					
cooking	3871.85	3937.49	1276.43	296.00	1084.80
idle	2470.21	2071.30	152.74	150.47	572.87
sleeping	1400.06	7989.06	231.36	559.09	1661.87
watching_tv	4456.64	4428.91	1725.52	6092.40	1295.40


```
In [15]: df_subset2.plot(kind='bar')
plt.title('Energy Consumption per appliance wrt activity in a year')
plt.xlabel('Activity')
plt.ylabel('Energy in kWh')
plt.savefig(FIG_PATH+'\Energy_Per_Device_Activity_Wise.png')
plt.show()
```



```
In [16]: df_subset3 = df_initial.pivot_table(index='user_present', columns='device_type', values='energy_kWh', aggfunc=df_subset3.round(2))
```

```
Out [16]: device_type  air_conditioner      fridge      light       tv      washer
user_present
0    803.07        4821.68     153.43    350.79    700.70
1   12198.76      18426.76   3386.05   7097.96   4614.94
```

```
In [17]: df_subset3.plot(kind='bar')
plt.title('Energy Consumption per appliance wrt presence of user in a year')
plt.xlabel('user present')
plt.ylabel('Energy in kWh')
plt.savefig(FIG_PATH+'\Energy_Per_Device_User_Presence_Wise.png')
plt.show()
```



```
In [18]: df_subset4 = df_initial.pivot_table(index='status', columns='device_type', values='energy_kwh', aggfunc='sum')
df_subset4.round(2)
```

```
Out [18]: device_type  air_conditioner    fridge    light      tv    washer
           status
           off      0.00        NaN      0.00      0.00      0.00
           on     13001.83   23248.45  3539.48  7448.75  5315.65
```


there is no device that consumes energy when off

```
In [19]: df_subset5 = df_initial.pivot_table(index='hour_of_day', columns='device_type', values='energy_kwh', aggfunc='sum')
df_subset5.round(2)
```

```
Out [19]: device_type  air_conditioner    fridge    light      tv    washer
           hour_of_day
           0      158.53   1000.05  97.34   70.82   139.99
           1      142.36   1002.17  30.89   69.48   141.92
           2      153.64   1000.74  30.10   69.47   126.74
           3      152.70   1005.37  28.04   71.74   139.75
           4      144.41   996.14   29.27   71.70   142.34
           5      138.14   999.31   68.53   69.61   145.22
           6      141.96   1000.57  233.94  71.42   387.80
           7      182.06   1001.42  168.40  74.26   387.77
           8      275.64   994.88   30.90  156.07   246.50
           9      383.66   985.69   28.23  287.43   140.34
           10     524.43   967.43   29.19  322.85   136.31
           11     662.08   946.19   30.85  317.42   134.69
           12     748.38   933.28   28.89  314.42   211.03
           13     789.10   912.90   29.57  320.85   213.12
           14     827.75   906.99   29.59  315.22   136.45
           15     826.02   908.66   31.62  317.78   135.28
           16     1005.22   909.43   28.46  240.00   126.80
           17     1328.95   921.83   193.86  108.98   134.52
           18     1312.26   940.58   471.00  188.58   134.84
           19     1113.84   956.73   559.65  818.14   386.61
           20     831.58   974.24   561.28  1282.21  389.58
           21     599.84   987.28   510.09  1206.63  392.10
           22     356.28   994.07   245.03  576.26   393.78
           23     203.01   1002.50  44.75   107.41  392.17
```

```
In [20]: df_subset5.plot(kind='line')
plt.title('Energy Consumption per appliance wrt month in a year')
plt.xlabel('Month')
plt.ylabel('Energy in kwh')
```

```
plt.savefig(FIG_PATH+'\Energy_Per_Device_Month_Wise.png')
plt.show()
```



```
In [21]: df_subset5 = df_initial.pivot_table(index='room', columns='device_type', values='energy_kwh', aggfunc='sum')
df_subset5.round(2)
```

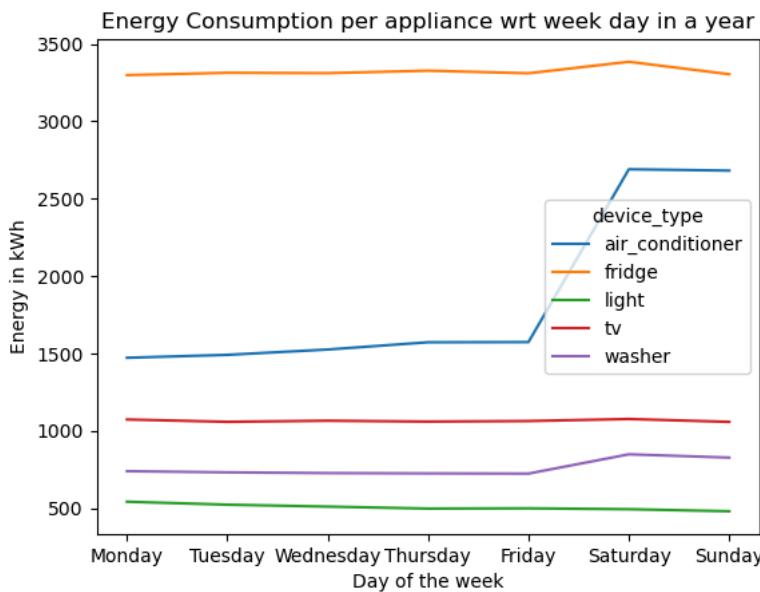
```
Out [21]: device_type  air_conditioner      fridge      light        tv      washer
           room
bedroom    13001.83       NaN       NaN       NaN       NaN
kitchen     NaN        23248.45       NaN       NaN       NaN
laundry_room   NaN       NaN       NaN       NaN      5315.65
living_room    NaN       NaN      3539.48    7448.75       NaN
```

from the above table it can be inferred that air_conditioner is installed in bedrooms only, fridge in kitchen, washer in laundry room while both light and tv are in living room

```
In [22]: df_subset6 = df_initial.pivot_table(index='day_of_week', columns='device_type', values='energy_kwh', aggfunc='sum')
weekdays={0: 'Monday',
1: 'Tuesday',
2: 'Wednesday',
3: 'Thursday',
4: 'Friday',
5: 'Saturday',
6: 'Sunday'}
df_subset6.index = df_subset6.index.map(weekdays)
df_subset6.round(2)
```

```
Out [22]: device_type  air_conditioner      fridge      light        tv      washer
           day_of_week
Monday     1471.41      3298.33     541.05    1072.53    738.38
Tuesday    1490.25      3313.36     521.87    1057.42    730.77
Wednesday   1524.76      3311.07     509.49    1064.57    726.06
Thursday    1571.67      3326.80     496.72    1058.90    724.06
Friday      1572.81      3310.20     498.31    1062.52    722.74
Saturday    2689.86      3384.52     492.65    1075.67    847.47
Sunday      2681.08      3304.16     479.40    1057.14    826.18
```


```
In [23]: df_subset6.plot(kind='line')
plt.title('Energy Consumption per appliance wrt week day in a year')
plt.xlabel('Day of the week')
plt.ylabel('Energy in kWh')
plt.savefig(FIG_PATH+'\Energy_Per_Device_WeekDay_Wise.png')
plt.show()
```



```
In [24]: df_subset7=df_initial.pivot_table(index='home_id', columns='device_type', values='energy_kWh', aggfunc='sum')
df_subset7.round(2)
```

```
Out [24]: device_type  air_conditioner    fridge     light      tv    washer
           home_id
1          1857.98    2328.13   363.49  1546.32   570.31
2          1850.95    2329.41   360.09  1552.05   588.86
3          1146.17    2327.58   353.87  544.56   443.84
4          1188.48    2324.04   346.71  544.80   394.11
5          1179.49    2328.48   347.30  546.20   412.70
6          1137.59    2322.11   359.20  540.63   585.30
7          1130.16    2321.25   365.72  539.04   560.68
8          1220.96    2320.35   342.18  551.20   584.50
9          1160.26    2323.73   352.99  548.24   595.07
10         1129.78    2323.37   347.95  535.70   580.27
```

```
In [25]: df_subset7.plot(kind='bar')
plt.title('Energy Consumption per appliance wrt household in a year')
plt.xlabel('Home ID')
plt.ylabel('Energy in kWh')
plt.savefig(FIG_PATH+'\Energy_Per_Device_Household_Wise.png')
plt.show()
```



```
In [26]: ## observing relation between values of 'price_kWh' and other features
```

```
In [27]: df_initial['price_kWh'].describe()
```

```
Out [27]: count    1.752000e+06
mean     2.250000e+03
std      5.590172e+02
```

```
min      1.500000e+03
25%     1.500000e+03
50%     2.500000e+03
75%     2.500000e+03
max      3.000000e+03
Name: price_kWh, dtype: float64
```

```
In [28]: for attr in df_initial.columns:
    if attr != 'energy_kWh' and (df_initial[attr].dtype != object):
        print(df_initial[[attr, 'energy_kWh']].corr())
        print()
```

```
          home_id  energy_kWh
home_id      1.000000   -0.046145
energy_kWh   -0.046145    1.000000

          timestamp  energy_kWh
timestamp    1.000000    0.008793
energy_kWh    0.008793    1.000000

          power_watt  energy_kWh
power_watt     1.0          1.0
energy_kWh     1.0          1.0

          user_present  energy_kWh
user_present   1.000000   0.108452
energy_kWh     0.108452   1.000000

          indoor_temp  energy_kWh
indoor_temp    1.000000   0.136976
energy_kWh     0.136976   1.000000

          outdoor_temp  energy_kWh
outdoor_temp   1.000000   0.136327
energy_kWh     0.136327   1.000000

          humidity  energy_kWh
humidity      1.000000   -0.044761
energy_kWh    -0.044761   1.000000

          light_level  energy_kWh
light_level    1.000000   0.019396
energy_kWh     0.019396   1.000000

          day_of_week  energy_kWh
day_of_week    1.000000   0.031299
energy_kWh     0.031299   1.000000

          hour_of_day  energy_kWh
hour_of_day    1.000000   0.145333
energy_kWh     0.145333   1.000000


          price_kWh  energy_kWh
price_kWh     1.000000   0.130418
energy_kWh     0.130418   1.000000
```

```
In [29]: df_subset8 = df_initial.pivot_table(index='hour_of_day', values='price_kWh', aggfunc='sum')
df_subset8.round(2)
```

```
Out [29]: price_kWh
```

hour_of_day	price_kWh
0	109500000
1	109500000
2	109500000
3	109500000
4	109500000
5	109500000
6	182500000
7	182500000
8	182500000
9	182500000
10	182500000
11	182500000
12	182500000
13	182500000
14	182500000
15	182500000
16	182500000
17	219000000
18	219000000
19	219000000
20	219000000
21	182500000
22	109500000
23	109500000

```
In [30]: df_subset8.plot(kind='line')
plt.title('Price variation through the day summed over year')
plt.xlabel('Hour of day')
plt.ylabel('Price in VND')
plt.savefig(FIG_PATH+'\VND_price_Distribution.png')
plt.show()
```



```
In [31]: df_subset9 = df_initial.pivot_table(index='home_id', values='price_kWh', aggfunc='sum')
df_subset9
```

```
Out [31]:      price_kWh
home_id
 1  394200000
 2  394200000
 3  394200000
 4  394200000
 5  394200000
 6  394200000
 7  394200000
 8  394200000
 9  394200000
10  394200000
```

the original dataset follows VND currency so these values might make sense by that country standards

Checking imbalance in record entries

```
In [32]: print(df_initial.columns)

Index(['home_id', 'timestamp', 'device_id', 'device_type', 'room', 'status',
       'power_watt', 'user_present', 'activity', 'indoor_temp', 'outdoor_temp',
       'humidity', 'light_level', 'day_of_week', 'hour_of_day', 'price_kWh',
       'energy_kwh'],
      dtype='object')
```

since the timestamp-appliance follow a panel-logging system, there will be balanced entries for the columns: timestamp, home_id, device_id, device_type, room, day_of_week and hour_of_day


for status:

```
In [33]: status_count=df_initial['status'].value_counts()
status_count
```

```
Out [33]: status
off    1151365
on     600635
Name: count, dtype: int64
```


```
In [34]: status_count.plot(kind='bar')
plt.title('Status Record Count')
plt.xlabel('Status')
```

```
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\Status_Record_Value_Count.png')
plt.show()
```



appliance-wise status imbalance

```
In [35]: pd.crosstab(df_initial['device_type'], df_initial['status']).plot(kind='bar', stacked=True)
plt.title('Device Type Count')
plt.xlabel('Device Type')
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\Device_Type_Record_Value_Count.png')
plt.show()
```




user_present record distribution:

```
In [36]: user_presence_count = df_initial['user_present'].value_counts()
user_presence_count
```


```
Out [36]: user_present
1    1378220
0     373780
Name: count, dtype: int64
```

```
In [38]: user_presence_count.plot(kind='bar')
plt.title('Count of user presence value records in dataset')
plt.xlabel('user presnt')
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\User_Presence_Record_Value_Count.png')
plt.show()
```



user_presence count home_id wise:


```
In [39]: pd.crosstab(df_initial['home_id'], df_initial['user_present']).plot(kind='bar', stacked=True)
plt.title('User presence Count per Household')
plt.xlabel('Home ID')
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\Household_User_Presence_Record_Value_Count.png')
plt.show()
```




```
In [40]: activity_count = df_initial['activity'].value_counts()
activity_count
```

```
Out [40]: activity
sleeping      582955
away          373780
watching_tv   335850
cooking        300395
idle           159020
Name: count, dtype: int64
```


```
In [41]: activity_count.plot(kind='bar')
plt.title('Count of activity records in dataset')
plt.xlabel('Activity')
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\Activity_Record_Value_Count.png')
plt.show()
```



```
In [42]: pd.crosstab(df_initial['home_id'], df_initial['activity']).plot(kind='bar', stacked=True)
plt.title('Activity Count per Household')
plt.xlabel('Home ID')
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\Activity_Household_Record_Value_Count.png')
plt.show()
```



```
In [43]: df_initial['energy_kWh'].plot(kind='hist')
plt.title('Energy Consumption value distribution')
plt.xlabel('Energy in kWh')
plt.ylabel('Frequency')
plt.savefig(FIG_PATH+'\Energy_Consumption_Distribution_Initial.png')
plt.show()
```



Most 15-minute intervals show very low energy use, while only a few intervals spike because of heavy appliances like ACs, washers, and fridges.

This kind of skewed pattern is normal in smart-home data and will be taken care of during normalization in later stages.

```
In [44]: df_initial['energy_kwh'].describe().round(2)
```

```
Out [44]: count    1752000.00
mean      0.03
std       0.06
min      0.00
25%      0.00
50%      0.00
75%      0.05
max      0.87
Name: energy_kwh, dtype: float64
```

```
In [ ]:
```